

The Kernel Trick, Gram Matrices, and Feature Extraction

CS6787 Lecture 4 — Fall 2017

Momentum for Principle Component Analysis

CS6787 Lecture 3.1 — Fall 2017

Principle Component Analysis

- Setting: find the dominant eigenvalue-eigenvector pair of a positive semidefinite symmetric matrix \mathbf{A} .

$$u_1 = \arg \max_x \frac{x^T A x}{x^T x}$$

- Many ways to write this problem, e.g.

$\|B\|_F$ is *Frobenius norm*

$$\sqrt{\lambda_1} u_1 = \arg \min_x \|xx^T - A\|_F^2$$

$$\|B\|_F^2 = \sum_i \sum_j B_{i,j}^2$$

PCA: A Non-Convex Problem

- PCA is **not convex** in any of these formulations
- **Why?** Think about the solutions to the problem: \mathbf{u} and $-\mathbf{u}$
 - Two distinct solutions \rightarrow can't be convex
- Can we still use momentum to run PCA more quickly?

Power Iteration

- Before we apply momentum, we need to choose what base algorithm we're using.
- Simplest algorithm: **power iteration**
 - Repeatedly multiply by the matrix A to get an answer

$$x_{t+1} = Ax_t$$

Why does Power Iteration Work?

- Let eigendecomposition of A be $A = \sum_{i=1}^n \lambda_i u_i u_i^T$
 - For $\lambda_1 > \lambda_2 \geq \lambda_3 \geq \dots \geq \lambda_n$
- PI converges **in direction** because cosine-squared of angle to \mathbf{u}_1 is

$$\cos^2(\theta) = \frac{(u_1^T x_t)^2}{\|x_t\|^2} = \frac{(u_1^T A^t x_0)^2}{\|A^t x_0\|^2}$$

What about a more general algorithm?

- Use both current iterate, and **history of past iterations**

$$x_{t+1} = \alpha_t A x_t + \beta_{t,1} x_{t-1} + \beta_{t,2} x_{t-2} + \cdots + \beta_{t,t} x_0$$

- for fixed parameters α and β
- **What class of functions can we express in this form?**
- Notice: x_t is always a degree- t polynomial in A times x_0
 - Can prove by induction that we can express **ANY polynomial**

Power Iteration and Polynomials

- Can also think of power iteration as a degree- t polynomial of \mathbf{A}

$$x_t = A^t x_0$$

- Is there a better degree- t polynomial to use than $f_t(x) = x^t$?
 - If we use a different polynomial, then we get

$$x_t = f_t(A)x_0 = \sum_{i=1}^n f_t(\lambda_i) u_i u_i^T x_0$$

- Ideal solution: choose polynomial with zeros at all non-dominant eigenvalues

Chebyshev Polynomials Again


- It turns out that Chebyshev polynomials solve this problem.
- Recall: $T_0(x) = 1$, $T_1(x) = x$ and

$$T_{n+1}(x) = 2xT_n(x) - T_{n-1}(x)$$

- Nice properties:


$$|x| \leq 1 \Rightarrow |T_n(x)| \leq 1$$

Chebyshev Polynomials




$$T_0(u) = 1$$

Chebyshev Polynomials




$$T_1(u) = u$$

Chebyshev Polynomials




$$T_2(u) = 2u^2 - 1$$


Chebyshev Polynomials




Chebyshev Polynomials



Chebyshev Polynomials



Chebyshev Polynomials



Chebyshev Polynomials Again

- It turns out that Chebyshev polynomials solve this problem.
- Recall: $T_0(x) = 0$, $T_1(x) = x$ and

$$T_{n+1}(x) = 2xT_n(x) - T_{n-1}(x)$$

- Nice properties:

$$|x| \leq 1 \Rightarrow |T_n(x)| \leq 1 \qquad T_n(1 + \epsilon) \approx \Theta \left(\left(1 + \sqrt{2\epsilon}\right)^n \right)$$

Using Chebyshev Polynomials

- So we can choose our polynomial f in terms of T
 - Want: $f_t(\lambda_1)$ to be as large as possible, subject to $|f_t(\lambda)| < 1$ for all $|\lambda| < \lambda_2$
- To make this work, set

$$f_n(x) = T_n \left(\frac{x}{\lambda_2} \right)$$

Convergence of Momentum PCA

$$\frac{x_t}{\|x_t\|} = \frac{\sum_{i=1}^n T_t \left(\frac{\lambda_i}{\lambda_2} \right) u_i u_i^T x_0}{\sqrt{\sum_{i=1}^n T_t^2 \left(\frac{\lambda_i}{\lambda_2} \right) (u_i^T x_0)^2}}$$

- Cosine-squared of angle to dominant component:

$$\cos^2(\theta) = \frac{(u_1^T x_t)^2}{\|x_t\|^2} = \frac{T_t^2 \left(\frac{\lambda_1}{\lambda_2} \right) (u_1^T x_0)^2}{\sum_{i=1}^n T_t^2 \left(\frac{\lambda_i}{\lambda_2} \right) (u_i^T x_0)^2}$$

Convergence of Momentum PCA (continued)

$$\begin{aligned}\cos^2(\theta) &= \frac{(u_1^T x_t)^2}{\|x_t\|^2} = \frac{T_t^2 \left(\frac{\lambda_1}{\lambda_2}\right) (u_1^T x_0)^2}{\sum_{i=1}^n T_t^2 \left(\frac{\lambda_i}{\lambda_2}\right) (u_i^T x_0)^2} \\ &= 1 - \frac{\sum_{i=2}^n T_t^2 \left(\frac{\lambda_i}{\lambda_2}\right) (u_i^T x_0)^2}{\sum_{i=1}^n T_t^2 \left(\frac{\lambda_i}{\lambda_2}\right) (u_i^T x_0)^2} \\ &\geq 1 - \frac{\sum_{i=2}^n (u_i^T x_0)^2}{T_t^2 \left(\frac{\lambda_1}{\lambda_2}\right) (u_1^T x_0)^2} = 1 - \Omega \left(T_t^{-2} \left(\frac{\lambda_1}{\lambda_2} \right) \right)\end{aligned}$$

Convergence of Momentum PCA (continued)

$$\begin{aligned}\cos^2(\theta) &\geq 1 - \Omega \left(T_t^{-2} \left(\frac{\lambda_1}{\lambda_2} \right) \right) = 1 - \Omega \left(T_t^{-2} \left(1 + \frac{\lambda_1 - \lambda_2}{\lambda_2} \right) \right) \\ &= 1 - \Omega \left(\left(1 + \sqrt{2 \frac{\lambda_1 - \lambda_2}{\lambda_2}} \right)^{-2t} \right)\end{aligned}$$

- Recall that standard power iteration had:

$$\cos^2(\theta) = 1 - \Omega \left(\left(\frac{\lambda_2}{\lambda_1} \right)^{2t} \right) = 1 - \Omega \left(\left(1 + \frac{\lambda_1 - \lambda_2}{\lambda_2} \right)^{-2t} \right)$$

- So the momentum rate is asymptotically faster than power iteration

Questions?

The Kernel Trick, Gram Matrices, and Feature Extraction

CS6787 Lecture 4 — Fall 2017

Basic Linear Models

- For classification using model vector \mathbf{w}

$$\text{output} = \text{sign}(w^T x)$$

- Optimization methods vary; here's logistic regression ($y_i \in \{-1, 1\}$)

$$\text{minimize}_w \frac{1}{n} \sum_{i=1}^n \log (1 + \exp(-w^T x_i y_i))$$

Benefits of Linear Models

- **Fast classification:** just one dot product
- **Fast training/learning:** just a few basic linear algebra operations
- **Drawback: limited expressivity**
 - Can only capture linear classification boundaries → bad for many problems
- How do we let linear models **represent a broader class of decision boundaries**, while **retaining the systems benefits**?

The Kernel Method

- Idea: in a linear model we can think about the **similarity** between two training examples \mathbf{x} and \mathbf{y} as being

$$x^T y$$

- This is related to the rate at which a random classifier will separate \mathbf{x} and \mathbf{y}
- Kernel methods replace this dot-product similarity with an arbitrary **Kernel function** that computes the similarity between \mathbf{x} and \mathbf{y}

$$K(x, y) : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$$

Kernel Properties

- **What properties do kernels need to have to be useful for learning?**
- Key property: kernel must be **symmetric** $K(x, y) = K(y, x)$
- Key property: kernel must be **positive semi-definite**

$$\forall c_i \in \mathbb{R}, x_i \in \mathcal{X}, \sum_{i=1}^n \sum_{j=1}^n c_i c_j K(x_i, x_j) \geq 0$$

- Can check that the dot product has this property

Facts about Positive Semidefinite Kernels

- Sum of two PSD kernels is a PSD kernel

$$K(x, y) = K_1(x, y) + K_2(x, y) \text{ is a PSD kernel}$$

- Product of two PSD kernels is a PSD kernel

$$K(x, y) = K_1(x, y)K_2(x, y) \text{ is a PSD kernel}$$

- Scaling by any function on both sides is a kernel

$$K(x, y) = f(x)K_1(x, y)f(y) \text{ is a PSD kernel}$$

Other Kernel Properties

- Useful property: kernels are often **non-negative**

$$K(x, y) \geq 0$$

- Useful property: kernels are often **scaled** such that





$$K(x, y) \leq 1, \text{ and } K(x, y) = 1 \Leftrightarrow x = y$$

- These properties capture the idea that the kernel is expressing the similarity between \mathbf{x} and \mathbf{y}

Common Kernels

- **Gaussian kernel/RBF kernel:** de-facto kernel in machine learning

$$K(x, y) = \exp(-\gamma \|x - y\|^2)$$

- We can validate that this is a kernel
 - Symmetric? 
 - Positive semi-definite?  **WHY?**
 - Non-negative? 
 - Scaled so that $\mathbf{K}(\mathbf{x}, \mathbf{x}) = 1$? 

Common Kernels (continued)

- **Linear kernel:** just the inner product $K(x, y) = x^T y$
- **Polynomial kernel:** $K(x, y) = (1 + x^T y)^p$
- **Laplacian kernel:** $K(x, y) = \exp(-\beta \|x - y\|_1)$
- Last layer of a neural network:
if last layer outputs $\phi(x)$, then kernel is $K(x, y) = \phi(x)^T \phi(y)$

Classifying with Kernels

- An equivalent way of writing a linear model on a training set is

$$\text{output}(x) = \text{sign} \left(\left(\sum_{i=1}^n w_i x_i \right)^T x \right)$$

- We can kernel-ize this by replacing the dot products with kernel evals

$$\text{output}(x) = \text{sign} \left(\sum_{i=1}^n w_i K(x_i, x) \right)$$

Learning with Kernels

- An equivalent way of writing linear-model logistic regression is

$$\text{minimize}_w \frac{1}{n} \sum_{i=1}^n \log \left(1 + \exp \left(- \left(\sum_{j=1}^n w_j x_j \right)^T x_i y_i \right) \right)$$

- We can kernel-ize this by replacing the dot products with kernel evals

$$\text{minimize}_w \frac{1}{n} \sum_{i=1}^n \log \left(1 + \exp \left(- \sum_{j=1}^n w_j y_i K(x_j, x_i) \right) \right)$$

The Computational Cost of Kernels

- Recall: benefit of learning with kernels is that **we can express a wider class of classification functions**
- Recall: another benefit is **linear classifier learning problems are “easy”** to solve because they are convex, and gradients easy to compute
- **Major cost of learning naively with Kernels:** have to evaluate $\mathbf{K}(\mathbf{x}, \mathbf{y})$
 - For SGD, need to do this effectively n times per update
 - Computationally intractable unless \mathbf{K} is very simple

The Gram Matrix

- Address this computational problem by **pre-computing the kernel function** for all pairs of training examples in the dataset.

$$G_{i,j} = K(x_i, x_j)$$

- Transforms the learning problem into

$$\text{minimize}_w \frac{1}{n} \sum_{i=1}^n \log (1 + \exp (-y_i e_i^T G w))$$

- This is much easier than recomputing the kernel at each iteration

Problems with the Gram Matrix

- Suppose we have n examples in our training set.
- **How much memory** is required to store the Gram matrix \mathbf{G} ?
- **What is the cost** of taking the product $\mathbf{G}_i \mathbf{w}$ to compute a gradient?
- What happens if we have **one hundred million training examples**?

Feature Extraction

- Simple case: let's imagine that \mathbf{X} is a finite set $\{1, 2, \dots, k\}$
- We can define our kernel as a matrix $M \in \mathbb{R}^{k \times k}$

$$M_{i,j} = K(i, j)$$

- Since M is positive semidefinite, it has a square root $U^T U = M$

$$\sum_{i=1}^k U_{k,i} U_{k,j} = M_{i,j} = K(i, j)$$

Feature Extraction (continued)

- So if we define a **feature mapping** $\phi(i) = Ue_i$ then

$$\phi(i)^T \phi(j) = \sum_{k=1}^k U_{k,i} U_{k,j} = M_{i,j} = K(i, j)$$

- The kernel is **equivalent to a dot product** in some space
- In fact, this is **true for all kernels**, not just finite ones

Classifying with feature maps

- Suppose that we can find a finite-dimensional feature map that satisfies

$$\phi(i)^T \phi(j) = K(i, j)$$

- Then we can simplify our classifier to

$$\begin{aligned} \text{output}(x) &= \text{sign} \left(\sum_{i=1}^n w_i K(x_i, x) \right) \\ &= \text{sign} \left(\sum_{i=1}^n w_i \phi(x_i)^T \phi(x) \right) = \text{sign} (u^T \phi(x)) \end{aligned}$$

Learning with feature maps

- Similarly we can simplify our learning objective to

$$\text{minimize}_u \frac{1}{n} \sum_{i=1}^n \log (1 + \exp (-u^T \phi(x_i) y_i))$$

- Take-away: this is just **transforming the input data, then running a linear classifier in the transformed space!**
- Computationally: **super efficient**
 - As long as we can transform and store the input data in an efficient way

Problems with Feature Maps

- The dimension of the transformed data may be **much larger than the dimension of the original data**.
- Suppose that the feature map is $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^D$ and there are \mathbf{n} examples
- **How much memory is needed** to store the transformed features?
- **What is the cost** of taking the product $u^T \phi(x_i)$ to compute a gradient?

Feature Maps vs. Gram Matrices

- **Systems trade-offs exist here.**
- When number of examples gets very large, **feature maps are better.**
- When transformed feature vectors have high dimensionality, **Gram matrices are better.**

Another Problem with Feature Maps

- Recall: I said there was always a feature map for any kernel such that

$$\phi(i)^T \phi(j) = K(i, j)$$

- But this feature map is **not always finite-dimensional**
 - For example, the Gaussian/RBF kernel has an infinite-dimensional feature map
 - **Many kernels we care about in ML have this property**
- What do we do if ϕ has infinite dimensions?
 - **We can't just compute with it normally!**

Solution: Approximate Feature Maps

- Find a finite-dimensional feature map so that

$$K(x, y) \approx \phi(x)^T \phi(y)$$

- Typically, we want to find a family of feature maps ϕ_t such that

$$\phi_D : \mathbb{R}^d \rightarrow \mathbb{R}^D$$

$$\lim_{D \rightarrow \infty} \phi_D(x)^T \phi_D(y) = K(x, y)$$

Types of Approximate Feature Maps

- Deterministic feature maps
 - Choose a fixed-a-priori method of approximating the kernel
 - Generally not very popular because of the way they scale with dimensions
- Random feature maps
 - Choose a feature map at random (typically each feature is independent) such that

$$\mathbf{E} [\phi(x)^T \phi(y)] = K(x, y)$$

- Then prove with high probability that over some region of interest

$$|\phi(x)^T \phi(y) - K(x, y)| \leq \epsilon$$

Types of Approximate Features (continued)

- Orthogonal randomized feature maps

- Intuition behind this: if we have a feature map where for some i and j

$$e_i^T \phi(x) \approx e_j^T \phi(x)$$

then we can't actually learn much from having both features.

- Strategy: choose the feature map at random, but subject to the constraint that the features be “orthogonal” in some way.

- Quasi-random feature maps

- Generate features using a low-discrepancy sequence rather than true randomness

Adaptive Feature Maps

- Everything before this **didn't take the data into account**
- Adaptive feature maps look at the actual training set and try to minimize the kernel approximation error using the training set as a guide
 - For example: we can do a random feature map, and then **fine-tune the randomness** to minimize the empirical error over the training set
 - Gaining in popularity
- Also, neural networks can be thought of as adaptive feature maps.

Systems Tradeoffs

- Lots of tradeoffs here
- Do we spend more work up-front constructing a more sophisticated approximation, to save work on learning algorithms?
- Would we rather scale with the data, or scale to more complicated problems?
- Another task for **metaparameter optimization**

Questions

- Upcoming things:
 - **Paper 2 review due tonight**
 - Paper 3 in class on Wednesday
 - Start thinking about the class project — it will come faster than you think!